Using `new`, `delete` and `delete[]`

Use `new` to create an object on the heap and get its address (`new` evaluates to the address of the newly-created object). Use `delete` on a heap object's address to deallocate it and free up memory.

ⓘ Using `delete` on a `nullptr` does nothing.

A dynamic object's lifetime begins when it's created via `new`, and its lifetime ends (and its dtor is called) only when you use `delete` on its address.

ⓘ Array elements are constructed from left-to-right (and are destroyed in the reverse order). Class members are initialized in the order that they're declared in the class body.

Dynamic arrays are created with `new` (which yields a pointer to its first element) and deleted by using `delete[]` *on the address of its first element*. (Note that you can't delete individual elements.)

```cpp
struct X {...}; class Y {...};

// default initialization
new int; // does nothing (atomic type)
new Y; // calls default Y ctor

// value initialization
int* ptr1 = new int(); // *ptr1 == 0
int* ptr2 = new int{}; // *ptr2 == 0
Y* ptr3 = new Y{}; // calls default ctor

// explicit initialization
int* ptr_d = new int(3);
X* ptr4 = new X{3}; // {} for struct
Y* ptr5 = new Y{3}; // {}, () for class
Y* ptr6 = new Y(3); // invokes Y ctor
```

```cpp
int N = 0;
cin >> N; // get num. elements from user
int* arr = new int[N]; // create size N array
arr[0] = 42; // sets arr's first element to 42
delete[] arr; // deletes entire dynamic array
```

ⓘ While dynamic arrays are still fixed-size, they're allowed to have their sizes determined at runtime, unlike non-dynamic arrays (which must have compile-time constant sizes).

❗ Deleting a dynamic array through a pointer with different static/dynamic types causes U.B.

⚠ If a non-dynamic class/struct object or array has dynamic sub-objects, the sub-objects can outlive the compound object they're part of.

	Storage Duration	Destructor Runs When:	Managed By:	Grows
Stack Memory	Local (automatic)	Scope `{}` ends	Compiler	✔ Yes
Heap Memory	Dynamic	Manually Deleted	Programmer	✔ Yes
Static Memory	Static	Program terminates	Compiler	✘ No

Memory Management Errors

The following mistakes cause undefined behavior:

1. *Use-after-free*: dereferencing a pointer after using `delete` on it (`delete` doesn't kill the pointer itself).
2. *Bad delete*: Using `delete` on a dynamic array or non-dynamic object, or using `delete[]` on anything other than the address of a dynamic array's first element.
3. *Double free*: Using `delete` on an address twice.

Forgetting to `delete` a heap object causes a **memory leak**, essentially "wasted" memory. Memory leaks are often caused by losing a dynamic object's address.

```cpp
int* func(int x) {
  int *y = new int(x);
  y = new int[x]; // Orphaned mem.
  return y;
} // This function leaks memory

int main() {
  int *a = func(5);
  int *b = a;
  delete b; // Wrong delete
  delete[] a; // double delete
  cout << a[2]; // use-after-free
} // We still have a memory leak...
```

RAII (Resource Aquisition Is Initialization)

RAII (or *scope-bound resource management*) is a strategy to prevent memory leaks by wrapping a dynamic object in a class. To create an RAII class:

1. Make the class's ctor allocate a dynamic object
2. Keep track of the object via a `private` pointer.
3. Make the class's dtor deallocate the object

This ties a dynamic object's lifetime to the scope of the non-dynamic class object that "holds" it.

```cpp
template <typename T>
class UnsortedSet {
private:
  T *arr; // ptr to underlying array
  int cap; // array's size limit
  int N; // current size of array
public:
  UnsortedSet() // ctor allocates arr
    : arr(new T[10]), cap(50), N(0) { }
  ~UnsortedSet() { delete[] arr; }
}; // dtor deallocates arr
```

⚠ The destructor for an RAII container only cleans up resources allocated for the container itself, not the memory used by dynamic objects that were stored in the container. If you have nested classes that manage dynamic memory, you need to implement custom destructors for all of them (and if you have a container of heap pointers, you can't expect the container to clean them up).

Shallow and Deep Copies

A *deep copy* of a class-type object is a copy whose data members are separate from the original's. A *shallow copy* is a copy whose data members are "shared" with the original object's. E.g. if you create a shallow copy of an `UnsortedSet`, both objects will have pointers to the *same* dynamic array. The concept only applies to classes that have "owning" pointers/references.

⚠ Compiler-generated copy constructors and assignment `=` operators create shallow copies, so you must implement custom versions of them for classes that have pointers to dynamic resources.

Ex: Copy Ctor Implementation	Ex: Assignment Operator Overload
| ```cpp
class IntSet { ...
// 1. Initialize copy's stack members
  IntSet(const IntSet& og)
    : cap(og.cap), N(og.size()) {
// 2. Create a new, separate dynamic array
    arr = new int[og.size()];
    for (int i = 0; i < og.size(); ++i) {
      arr[i] = og.arr[i];
    } // 3. Copy elements to new array
  }
};
``` | ```cpp
IntSet& operator=(const IntSet& og) {
  if (this == &og) { return *this; }
  delete[] this->elts;
  this->cap = og.cap;
  this->N = og.size();
  this->elts = new int[og.size()];
  for (int i = 0; i < og.size(); ++i) {
    this->elts[i] = og.elts[i];
  }
  return *this;
}
``` |

ⓘ You can explicitly call a copy constructor by simply passing a pre-existing class object to it. Also, if you initialize a new class object from an existing class object using the assignment operator, that actually makes an implicit call to the class's copy constructor.

Destructors and Polymorphism (Virtual Destructors)

If you try to delete a derived-class object by calling `delete` on a base-class pointer to the object, the base class needs to have a `virtual` destructor defined or else undefined behavior will occur.

The compiler-generated default destructor is not `virtual`, so it's best practice to define a custom virtual destructor for polymorphic classes (i.e., classes that make use of inheritance).

⚠ If you call `delete` on an object whose static type and dynamic type are different, the static type must be a base class of the dynamic type (or U.B. will occur even with a `virtual` destructor).

ⓘ You can prevent derived class objects from being deleted through base-class pointers by declaring the base class destructor `protected` and non-`virtual`.

The Rule of Three

A class's underline{destructor}, underline{copy constructor}, and underline{assignment operator} make up the *Big Three*.

Rule of Three: a class that needs a custom version of one of the Big Three usually needs custom versions of the other two (generally, this describes classes that manage dynamic memory).

ⓘ Not all classes with pointer data members need the Big Three. Ex: Linked List iterators store internal `Node*` pointers, but they don't need the Big Three because the List dtor is responsible for managing the nodes. Generally, classes that manage dynamic memory need the Big Three.

The Linked List Interface (Single and Double)

A *linked list* is a container that stores its elements non-contiguously in memory; its elements are represented as distinct "nodes" that are linked together via pointers. If nodes are linked in one direction, the list is singly-linked, and if nodes are linked in both directions, the list is doubly-linked.

ⓘ Since there's no need to maintain contiguity, inserting/removing nodes is O(1) at any position in a list, as the only adjustment you need to make is re-pointing the previous node's `next` pointer (and the following node's `prev` pointer in a doubly-linked list) instead of shifting every element. This does mean lists have O(n) instead of O(1) indexing, though.

Linked List Operations

Properties and Types of Recursion

A *recursive function* is a function that is defined in terms of itself (i.e., that calls itself). Recursive functions break problems down into two types of "cases":

- *Base cases*, problems that can be solved without recursion (and act as the end condition for a recursive loop). Ex: 0 and 1 for a factorial function.
- *Recursive cases*, which the function breaks down into simpler sub-problems and then passes to itself (until the problem has been reduced to base cases).

There are three types of recursive functions:

1. *Linear recursive*: functions that make no more than one call to themselves for each invocation.
2. *Tail recursive*: a type of linear recursive function whose final instruction is the recursive call.
3. *Tree recursive*: functions that can make *multiple* recursive calls to themselves in one stack frame (note: they're not limited to tree data structures).

⚠ A recursive call being on the last line of a function does NOT guarantee that the function is tail recursive—the recursive call must be the last *instruction* executed by the function.

Linear Recursion (Non-Tail)
Tree Recursion
| ```cpp
bool SameTree(TreeNode* p,TreeNode* q) {
  if (!p || !q) {return (!p == !q);}
  if (p->val != q->val) {return false;}
  return SameTree(p->L, q->L)
      && SameTree(p->R, q->R);
} // check if two trees are identical

void func(int n) {
  for (int i = 0; i < n; ++i) func(i);
} // Yes, this is tree-recursive
``` |

```cpp
int factorial(int n) {
  return (n <= 1) ? 1 : n * factorial(n - 1);
} /* This is NOT tail-recursive because the last instruction the function executes is
     multiplying the recursive result by n, not the recursive call itself. */
```

Memory Usage of Recursive Functions

Non-tail linear recursive functions allocate an additional stack frame with each recursive call. However, the compiler can optimize a tail recursive function to reuse the same stack frame, which means that tail-recursive functions can be optimized to use a constant number of stack frames.

⚠ The total number of recursive calls a function makes to itself does not indicate its space complexity, since not all stack frames have to exist at the same time. Instead, look at the underline{number of return statements until the base case}. For example, a recursive algorithm to traverse a balanced BST of height h uses a maximum of O(h) space—not O(2^h) space—at any given time.

Structural Recursion

Structural recursion is when an abstract data type is defined in terms of itself. The `Node` structs of a linked list are one example: each `Node` contains a pointer to another `Node`. Recursive structures have base cases and recursive cases just like recursive functions do: for a linked list, the "base case" is an empty list, and the recursive cases are non-empty lists.

Recursion vs Iteration (and Converting Between Them)

Binary Search Trees

A *binary tree* is a tree data structure where each *node* points to at most 2 children. A *binary search tree* is a binary tree where, for *any* node, every node in the left subtree has a lower value and every node in the right subtree has a greater value. (An empty tree is also a BST.)

A BST is *balanced* if the heights of the left and right subtrees at any node differ by no more than 1, so the height of a balanced BST with n nodes is roughly $\log_2(n)$. A *leaf* of a BST is a childless node.

BST Traversal

There are multiple ways to traverse a binary tree:

1. *Inorder traversal*: recursively process the *left* subtree → process the *head* (current) node → recursively process the *right* subtree.
2. *Preorder traversal*: process the *head* (current) node → recursively process the *left* subtree → recursively process the *right* subtree.
3. *Postorder traversal*: recursively process the *left* subtree → recursively process the *right* subtree. → process the *head*.

ⓘ Inorder traversal of a BST visits nodes by ascending value.

Inorder Traversal (LPR)

1	5	8	43	51	52	83	87

Preorder Traversal (PLR)

51	5	1	8	43	87	52	83

Postorder Traversal (LRP)

1	8	43	5	83	52	87	51

```cpp
void preOrder(Node* p) {
  if (!p) return;
  process(p->val);
  preOrder(p->left);
  preOrder(p->right);
} // preorder traversal
```

```cpp
void inOrder(Node* p) {
  if (!p) return;
  inOrder(p->left);
  process(p->val);
  inOrder(p->right);
} // inorder traversal
```

```cpp
void postOrder(Node* p) {
  if (!p) return;
  postOrder(p->left);
  postOrder(p->right);
  process(p->val);
} // postorder traversal
```

Recursive Big Three

Operations on BSTs vs Sorted/Unsorted Sets

Search: due to the sorting invariant, searching a balanced BST for a value takes O($\log n$) time.

Recursive Binary Search

```cpp
template <typename Iter, typename T>
int binarySearch(Iter left, Iter right, const T& val) {
  int size = (right - left);
  if (size <= 0) return -1;
  Iter mid = left + (size / 2);
  if (*mid == val) return (mid - left);
  if (*mid > val) return binarySearch(left, mid, val);
  return binarySearch(mid + 1, right, val);
} // eliminates half the search space each loop
```

ⓘ `std::map` and `std::set` are implemented using *self-balancing* binary search trees, which is how they maintain $\Theta(\log n)$ insertion, removal, and searching even in the worst case.

*Where n = number of nodes in the tree

Best case: $\Omega(n)$

Worst case: $O(n)$

Average case: $\Theta(n)$

Iterators and Functors

Functors (function objects): *first-class* entities that provide the same interface as a function. They can be created from classes that overload the function-call operator, i.e., `operator()`.

ⓘ First-class entities can *store state* (store and access information), be created at runtime, and be passed as an argument to/returned by a function.

ⓘ `operator()` can only be overloaded from within a class body. It and `operator[]` are also the only operators that can be overloaded as `static` member functions.

Function Pointers

Iterators are a type of object that are designed to emulate the interface of a pointer.

	std::sort(it1,it2)	it2 - it1	it[n]	it += n	it1 < it2	--it	++it	it1 == it2	*it
Random Access	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes
Bidirectional	✘ No	✘ No	✘ No	✘ No	✘ No	✔ Yes	✔ Yes	✔ Yes	✔ Yes
Forward	✘ No	✘ No	✘ No	✘ No	✘ No	✘ No	✔ Yes	✔ Yes	✔ Yes

Iterator	array<T,N>::iterator	vector<T>::iterator	list<T>::iterator	map<Key,T>::iterator	set<Key>::iterator
Iterator Class	Random Access	Random Access	Bidirectional	Bidirectional	Bidirectional
Container Resizing	N/A	All iterators invalidated	Dynamic	Empty container	Linked
Insertion	Sequential	Iterators at/above point of insertion are invalidated	Unaffected	Unaffected	Unaffected
Erasure	Sequential	Iterators at/above point of erasure are invalidated	Dynamic	Empty container	Linked

Friend Classes

Declaring a class `S` as a `friend` of class `T` makes the private data members of `S` accessible in the scope of `T` (but not vice versa— `T`'s private members don't become accessible to `S`).

```cpp
friend class F;
friend F;
{
```

Error Handling and Exceptions

ⓘ A `try` block can have multiple `catch` blocks associated with it.

Uncaught exceptions will make the program terminate.

C++ Stuff and Impostor Syndrome

Impostor syndrome is characterized by doubt in one's abilities that aren't backed up by evidence and an inability to recognize/take credit for one's achievements. These feelings are common and can affect anyone (although people from underrepresented groups may be more susceptible to it).

A *static member function* does *not* have access to the `this` pointer and also can't access non-`static` data members (reminder: `static` members are "shared" between all instances of a class).

Templates and function overloading are forms of *compile-time polymorphism*, while subtype polymorphism is a form of *runtime polymorphism*.

Container ADTs

`static` keyword: used to make one copy of a class data member "shared" between all instances of that class. A `static` data member has static storage duration but exists only within the scope of a class.

stack: a container that's designed to operate in a LIFO order.

queue: a container designed to operate in a first-in/first-out (FIFO) order.

Container	Interface Operations - Optimal Implementations are All O(1)				
Stack	empty	size	top (next to pop)		push_back	pop_back
Queue	empty	size	front (next)	back (last)	push_back	pop_front

ⓘ An efficient way to implement a queue is to create a vector with free space at both ends (a **ring/circular buffer**). This requires keeping track of the data's "head" (inclusive) and "tail" (exclusive).

Useful `std::vector` Functions and Constructors				
.size()	v[i]/.at(i)	.push_back(val)	.pop_back()	.resize(n)
.front()	.back()	.erase(iterator)	.empty()	.clear()
vector<T> v2(v1.begin(), v1.end())			vector<T> v2(v1)	

* These refer to the array-based sets we saw in class, *not* STL sets.

C++ Standard Library Containers

`std::array`: Containers with *compile-time constant* sizes that store elements in contiguous memory.

`std::vector`: resizable containers that store elements at the front and free space at the back.

⚠ Element insertions/deletions will invalidate `std::vector` iterators *at and after* the modified index. Also, any operation that changes the maximum capacity of a vector invalidates all existing iterators.

`std::list`: a container whose elements are linked via pointers (i.e., in non-contiguous memory).

`std::map`: an associative array that maps unique keys to values. Keys act like indexes for a map.

`std::set`: an associative sorted container that only stores unique keys.

ⓘ Removing an element from a `std::list`, `std::set`, or `std::map` only invalidates iterators to the removed element. Also, inserting an element doesn't ever invalidate iterators for those containers.

C++ Standard Library Containers

```cpp
std::array<int, 4> arr = {-4,0,3,6};
std::array arr2{1, 2}; // Only way to omit
    size
std::list<int> doubly_linked = {1,2,3};
std::forward_list<int> singly_linked = {4,5,6
    };
std::set<int> nums = {3,2,2,1}; // {1,2,3}

std::map<string, int> EECS = { {"Bill", 183}
    };
cout << EECS["Bill"] << endl; // prints 183
// Two ways to insert into a map:
EECS["Emily"] = 203;
EECS.insert(pair<string, int>("James", 280));
```

Useful `std::map` functions

```cpp
// Returns iterator to the pair with Key == k
// Returns .end() iterator if no such pair
    exists
iterator find(const Key_type& k) const;

// Inserts a <key, value> std::pair into a map
// Returns <iterator, false> if Key already in
    use
pair<iterator,bool> insert(const Pair_type&
    pair);

/* Finds or enters a value for a given key,
    then
returns a reference to the associated value */
Value_type& operator[](const Key_type& key);
```

Templates (Parametric Polymorphism)

Templates: special functions that take a data type as a parameter at compile time and instantiate an object or function compatible with that type. They help reduce code duplication in container interfaces.

Class Template Syntax

```cpp
template <typename T>
class UnsortedSet {
public:
  void insert(T my_val);
  bool contains(T my_val) const;
  int size() const;
private:
  T elts[ELTS_CAPACITY];
  int elts_size;
  ...
}; // Syntax: UnsortedSet<type> s;
```

Function Template Syntax

```cpp
// Note: "class" also works in place of "typename"
template <typename T> // "T" is also an arbitrary name
T maxValue(const T &valA, const T &valB) {
  return (valB > valA) ? valB : valA;
} // This function returns the greater of valA and valB
// Syntax to call it: maxValue<int/double/etc>(...);
```

Templated Class Member Function Syntax

```cpp
template <typename T> // Necessary if outside class body
void UnsortedSet<T>::insert(T my_val) { ··· }
```

Iterators, Traversal by Iterator and Range-Based Loops

Iterators: objects that have the same interface as pointers; they provide a general interface for traversing and accessing the elements of different types of containers. Note that each container's iterator has its own invalidation conditions (and that using an invalidated iterator causes undefined behavior).

ⓘ `std::begin()` returns an iterator to the start of an STL container. `std::end()` returns an iterator that's *1 past* the end of an STL container (the iterator returned by `std::end()` should not be dereferenced).

	std::sort(it1,it2)	it2 - it1	it[n]	it += n	<, <=	--	++	==, !=	*it
Random Access Iterators	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes	✔ Yes
Bidirectional Iterators	✘ No	✘ No	✘ No	✘ No	✘ No	✔ Yes	✔ Yes	✔ Yes	✔ Yes
Forward Iterators	✘ No	✘ No	✘ No	✘ No	✘ No	✘ No	✔ Yes	✔ Yes	✔ Yes

Traversal by iterator: a more general form of traversing a container data type by pointer.

```cpp
vector<int> v(3, -1); // this syntax initializes v to
    {-1,-1,-1}
for (vector<int>::iterator it = v.begin(); it != v.end(); ++it)
  { cout << *it << endl; } // ::const_iterator if const vector
```

Range-Based For Loop (Works on Any Sequence Traversible by Iterator)

```cpp
vector<int> v = { 1, 2, 3, 4 };
// for (<type> <variable> : <sequence>) {...}
for (int item : v) { // works with arrays too
  cout << item << endl;
} // could also declare item as const or a ref
```

```cpp
// Compiler translation of range-based for
    loop
for (auto it = v.begin(); it != v.end(); ++it)
    {
  int item = *it;
  cout << item << endl;
} // auto keyword makes compiler deduce type
```